WHAT IT IS

What will renting GPUs cost, and what if prices change?

WHO USES IT

AI teams that rent GPUs, and the people who budget for them.

THE QUESTION

We need 32 H100 GPUs for 720 hours (30 days). What will it cost, which region is lowest, and what if prices rise?

THE ANSWER

A cost from Azure's public price list, the range across regions, and how much to set aside for price rises.

Not included: storage, data transfer, IP addresses, licences, taxes, discounts, reserved-price plans, Spot interruptions and availability limits.

EXAMPLE · INVENTED PRICES

32 H100 GPUs for 720 hours

USD 276,480

at the lowest listed price, in US East

Across 5 regions:
USD 276,480 – 345,600

If prices rise 25%: + USD 69,120

If prices rise 50%: + USD 138,240

Scenario calculation, not a forecast.

WHY IT MATTERS FOR MARKETS

GPU time is becoming a commodity.

A BIG, MOVING COST

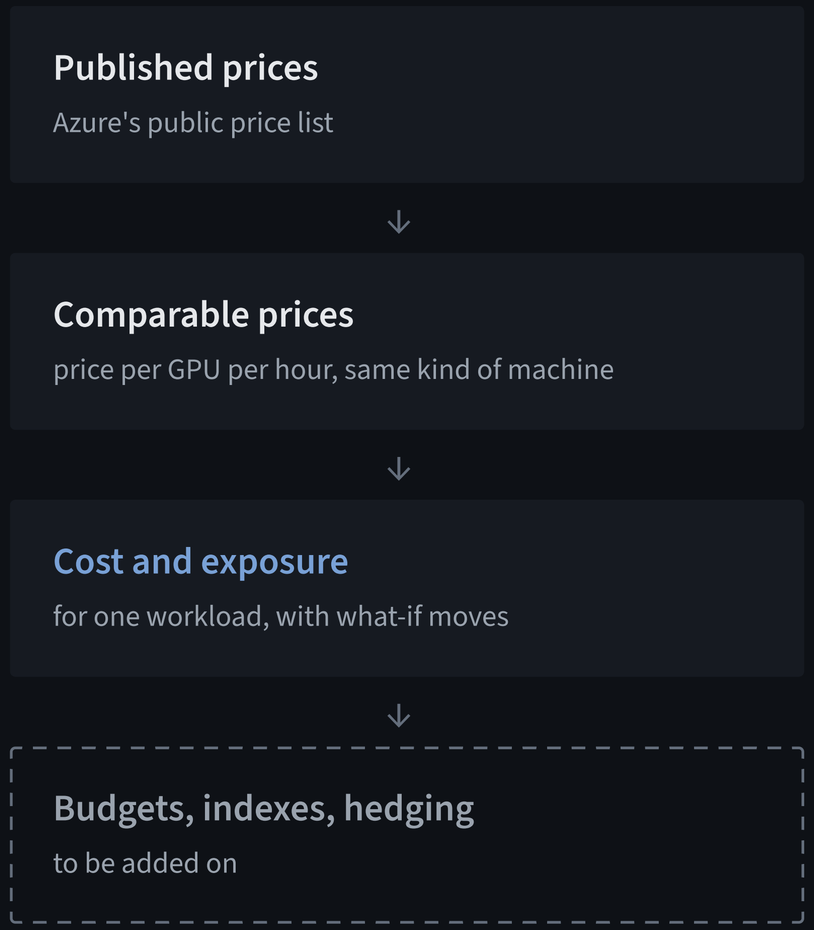
For AI companies, GPUs are one of the biggest bills, and prices vary by region and contract.

MARKETS ARE FORMING

Exchanges are preparing futures on GPU rental prices. CME's are under regulatory review; ICE's have no start date yet (October 2026).

THE GROUNDWORK

Budgeting, price indexes and hedging all need the same base: comparable prices and a clear measure of how exposed a budget is.



This project builds the first three steps. The last is to be added on.

HOW IT WORKS

Four simple steps

1

Collect

Download Azure's public GPU price list. Keep the original files, so every number can be traced.

2

Make comparable

Turn each price into a price per GPU per hour, and only compare the same kind of machine.

3

Price the workload

GPUs × hours × price. Extras like storage and tax are shown as not included.

4

Test price moves

Show the cost if prices move from −50% to +100%. What-ifs, not forecasts.

WORKED EXAMPLE · INVENTED PRICES

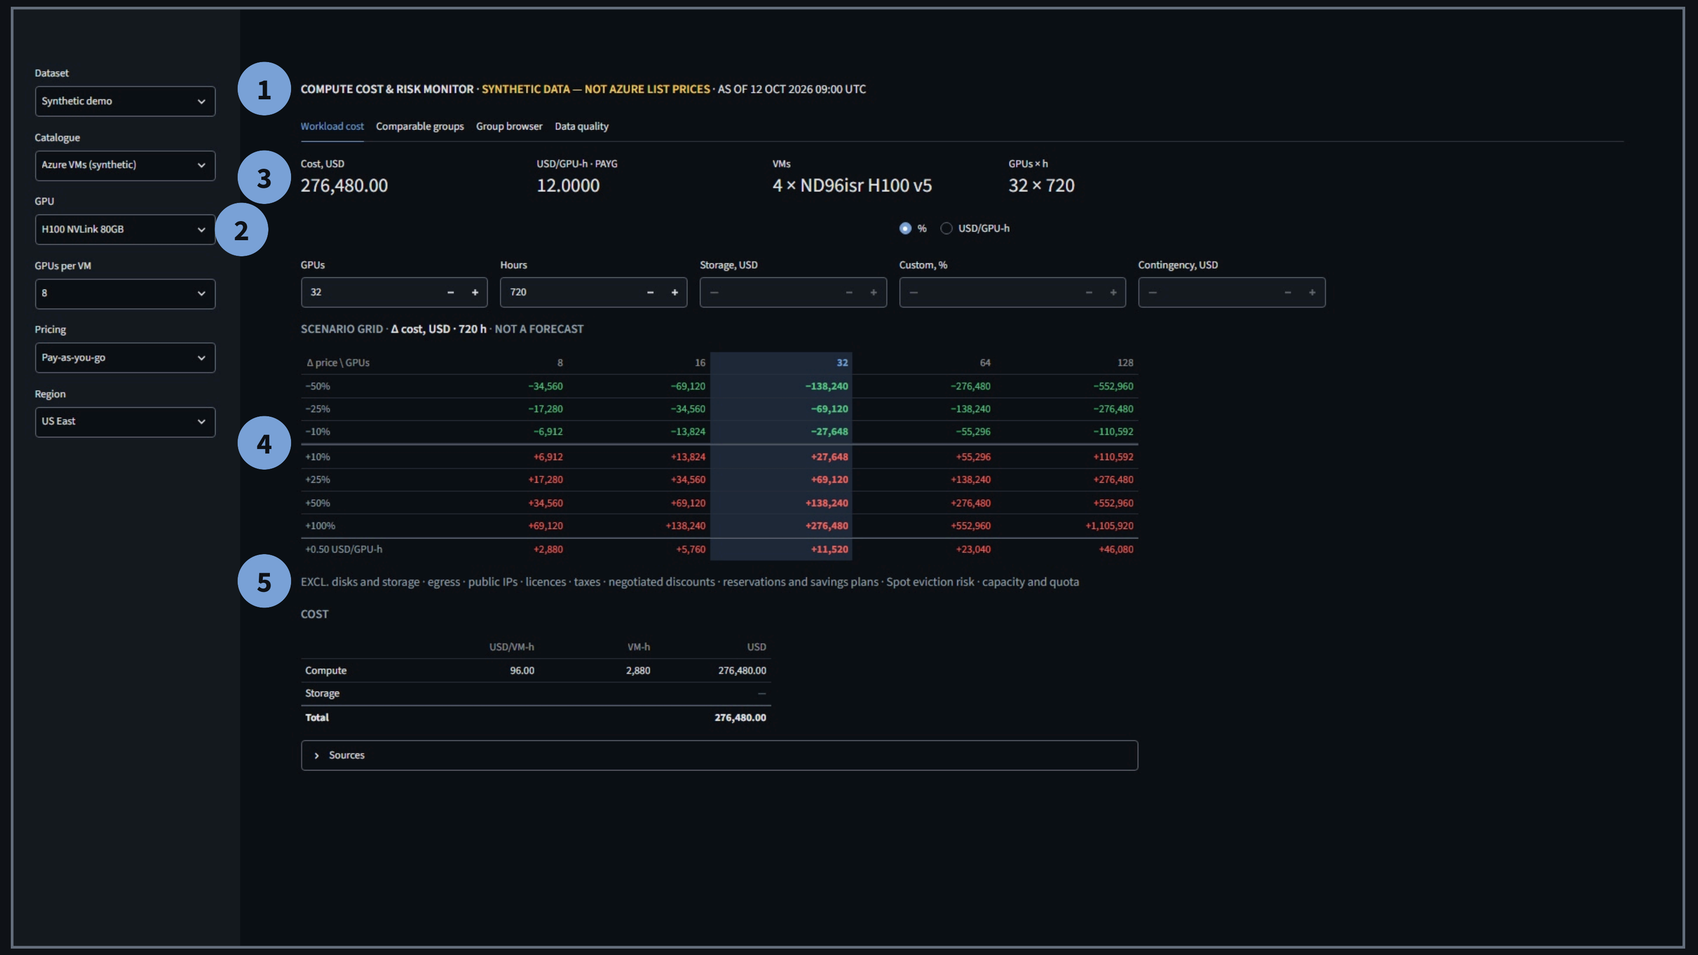
32 GPUs × 720 hours × USD 12.00 per GPU-hour = USD 276,480. If prices rise 25%: + USD 69,120.

Scenario calculation, not a forecast. Not included: storage, data transfer, IP addresses, licences, taxes, discounts, reserved-price plans, Spot interruptions and availability limits.

WHAT YOU SEE

The screen, part by part

1. Label: invented demo prices, and their date.
2. Your choices: GPU, machine size, pricing type, region.
3. The answer: total cost and price per GPU-hour.
4. What-if grid: extra cost in red, savings in green.
5. What's not included in the cost.



Demo with invented prices: 32 H100 GPUs, US East, 720 hours.

HOW IT USES AI

AI-powered market brief

1. Code calculates every number.
2. A second, independent check recalculates every number. If one disagrees, nothing is shown.
3. Claude (AI) writes the summary's words. It sees blanks like [[cost]], and code fills in the numbers.

FROM THE SUMMARY · INVENTED PRICES

“Running 32 H100 NVLink 80GB GPUs as 4 × ND96isr H100 v5 for 720 hours costs USD 276,480.00 at the lowest listed pay-as-you-go price, in US East at 12.0000 USD/GPU-h.”

LIMITS AND NEXT STEPS

What it is not, and what comes next

IT IS NOT

- A market price. It uses Azure's list prices, not what customers actually pay.
- A forecast. Price moves are what-ifs.
- Real prices in public. Microsoft's terms don't allow sharing them, so the demo uses invented ones.

LIMITS

- Little price history yet, so no trends.
- Only GPUs named in Microsoft's documentation are used; others are counted, not guessed.
- Risks and safeguards are written up using banking's model-risk standard (SR 11-7).

NEXT Collect more price history · add a second cloud provider